

Tools for Data Input/Output

HBD

Technology Map

“EXTRACT”
(Acquire)



“TRANSFORM”
(Process)



“LOAD”
(STORE)



APACHE
HBASE



Introduction

Applications implemented on **Hadoop** involve ingestion of disparate data types from **multiple sources** and with **differing requirements for frequency of ingestion**.

Common data sources for Hadoop include:

- » Traditional data management systems such as relational databases and mainframes
- » Logs, machine-generated data, and other forms of event data
- » Files being imported from existing enterprise data storage systems

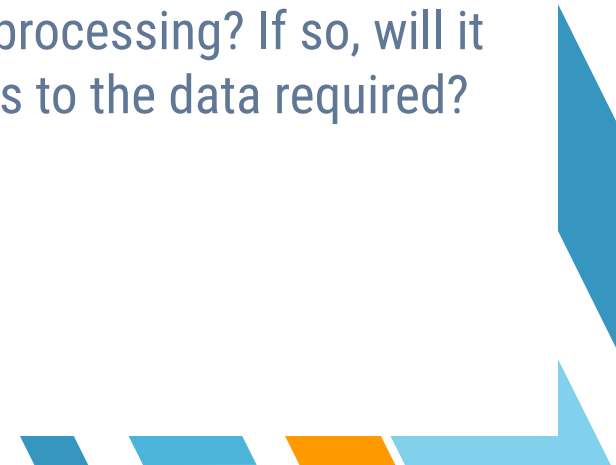
Hadoop (and other Big data tool) **start with the data in an HDFS** (or equivalent) so we need to **upload the data** to it.

Usually the source/destination systems outside the cluster are the bottleneck.





Considerations

- » **Timeliness of data ingestion and accessibility:** What are the requirements around how often data needs to be ingested? How soon does data need to be available to downstream processing?
 - » **Incremental updates:** How will new data be added? Does it need to be appended to existing data? Or overwrite existing data?
 - » **Data access and processing:** Will the data be used in processing? If so, will it be used in batch processing jobs? Or is random access to the data required?
- 

SQOOP

Command line tool designed to transfer data between Hadoop and relational databases.






Apache Sqoop


Usual workflow:

- » Import data from a relational database management system (RDBMS) into the Hadoop Distributed File System (HDFS) with Sqoop.
- » Transform the data with Hadoop MapReduce/Hive/Pig/etc.
- » Export data into an RDBMS with Sqoop.

Uses MapReduce to import and export the data, which provides parallel operation as well as fault tolerance.

Currently Sqoop 2 is available but not widely used and not feature complete. Sqoop 2 moves from shell client submit maps job to import, to a server with map reduce jobs and security. **We will talk about Sqoop 1.**



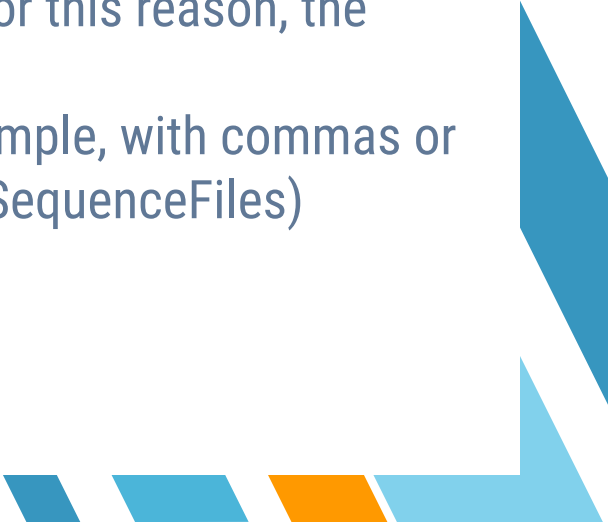


Sqoop Import

From database to HDFS



Apache Sqoop Import


- » The **input** to the import process is a **database table**.
 - » Sqoop will **read** the table **row-by-row into HDFS**.
 - » The **output** of this import process **is a set of files** containing a copy of the imported table.
 - » The import process **is performed in parallel**. For this reason, the output will be in **multiple files**.
 - » These files may be delimited **text files** (for example, with commas or tabs separating each field), **or binary** (Avro or SequenceFiles) containing serialized record data.
- 

Apache Sqoop Import Example

```
$. /sqoop import \  
--append --check-column "modified_date" --incremental "lastmodified" \  
--connect 'jdbc:mysql://db-host:3306/db' \  
--username my_sql_user --password-file '/user/hadoop/my_sql_user.password' \  
--table transactions --split-by "transaction_id" \  
--direct -m 8 \  
--hive-home "/home/hadoop/hive/" --hive-import \  
--hive-table "staging.transactions" --null-string '\\N' --null-non-string '\\N' \  
--hive-partition-key day --hive-partition-value ${current_date}
```



Apache Sqoop Import Parameters

- » **--connect:** jdbc connection string for the database.
 - » **--username:** username to connect to the database.
 - » **-P|--password-file|--password:** way to supply the password for the db connection.
 - » **--append:** Append data to an existing dataset in HDFS.
 - » **--direct:** use database specific driver (e.g. mysqldump for MySQL).
 - » **--columns <col,col,col...>:** Columns to import from table.
 - » **-m,--num-mappers <n>:** Use n map tasks to import in parallel.
 - » **-e,--query <statement>:** Import the results of statement.
 - » **--table <table-name>:** Table to read.
 - » **--where <where clause>:** WHERE clause to use during import.
 - » **--hive-import,--hive-overwrite:** additional options to import to Hive.
- 



Apache Sqoop Incremental Import

Retrieve **only rows newer** than some previously imported set of rows.


- » **--check-column (col)**: Specifies the column to be examined when determining which rows to import.
- » **--incremental (mode)**: Specifies how Sqoop determines which rows are new. Legal values for mode include **append** (if we only insert and we need to check by id) and **lastmodified** (if we update and we can check a last modified field).
- » **--last-value (value)**: Specifies the maximum value of the check column from the previous import (last id or last timestamp).

*At the end of an incremental import, the value which should be specified as **--last-value** for a subsequent import **is printed to the screen**.*





Apache Sqoop Import Internals

- » A Java class is generated during the import process, which can encapsulate one row of the imported table.
 - » This class is used during the import process.
 - » The Java source code for this class is also provided to you, for use in subsequent MapReduce processing of the data.
 - » This class can serialize and deserialize data to and from the SequenceFile format.
 - » It can also parse the delimited-text form of a record.
- 




Apache Sqoop import all tables in database


The following conditions must be met:

- » Each table must have a single-column primary key.
- » You must intend to import all columns of each table.
- » You must not intend to use non-default splitting column, nor impose any conditions via a WHERE clause.

```
$ sqoop import-all-tables --connect jdbc:mysql://db.foo.com/corp
```

After the import you would find in HDFS the following folders (on per table):

- » **/user/someuser/EMPLOYEES**
 - » **/user/someuser/PAYCHECKS**
 - » **/user/someuser/OFFICE_SUPPLIES**
- 



Sqoop Export

From HDFS to RDBMS




Apache Sqoop Export

Sqoop's export process will:

- » **Read** a set of delimited **text files from HDFS** in parallel.
- » **Parse** them into records.
- » **Insert** them as new rows in a **target database table**.

Considerations


- » The **target table must already exist** in the database.
 - » The **default operation** is to transform these into a set of **INSERT statements** that inject the records into the database.
 - » In "**update mode**", Sqoop will generate **UPDATE statements** that replace existing records in the database.
 - » In "**call mode**" Sqoop will make a **stored procedure call** for each record.
- 

Apache Sqoop Export Example

```
sqoop export \  
  -Dsqoop.export.records.per.statement=50 \  
  -Dsqoop.export.statements.per.transaction=50 \  
  --connect 'jdbc:mysql://db-host/db' \  
  --username my_sql_user --password 'pass123' \  
  --table transactions \  
  -m 12 \  
  --fields-terminated-by '\01' \  
  --input-null-string '\\N' --input-null-non-string '\\N' \  
  --export-dir /user/hive/warehouse/database.db/transactions \  
  -- --default-character-set=utf8mb4
```




Apache Sqoop Export arguments

- » **--export-dir <dir>**: HDFS source path for the export.
 - » **-m,--num-mappers <n>**: Use n map tasks to export in parallel.
 - » **--table <table-name>**: Table to populate.
 - » **--update-mode <mode>**: Specify how updates are performed when new rows are found with non-matching keys in database. Legal values for mode include **updateonly** (default) and **allowinsert**.
- 



Apache Sqoop Export staging table

Sqoop **breaks down export process into multiple transactions** which can mean partial data being committed if a job fails in the middle.

The **solution is specifying a staging table via the --staging-table** option which acts as an auxiliary table that is used to stage exported data. The staged data is finally moved to the destination table in a single transaction.

Considerations:

You must create **the staging table prior** to running the export job.

This table must be **structurally identical** to the target table.

This table should either **be empty** before the export job runs, or the **--clear-staging-table option** must be specified.





Apache Sqoop Export Consideration

By default **sqoop-export** appends new rows to a table using an **INSERT** statement.

The export process will fail if an **INSERT** statement fails.

If you specify the **--update-key** argument, Sqoop will instead **modify an existing dataset in the database**.

Each input record is treated as an **UPDATE** statement that modifies an existing row.


The row modified by a statement is determined by the column name(s) specified with **--update-key**.





Apache Sqoop Tools

```
$ sqoop <tool>
```

- » **codegen**: Generate code to interact with database records.
 - » **create-hive-table**: Import a table definition into Hive.
 - » **eval**: Evaluate a SQL statement and display the results.
 - » **export**: Export an HDFS directory to a database table.
 - » **import**: Import a table from a database to HDFS.
 - » **import-all-tables**: Import tables from a database to HDFS.
 - » **list-databases**: List available databases on a server.
 - » **list-tables**: List available tables in a database.
- 

Apache Sqoop Metastore

The **Sqoop Metastore** is a powerful part of Sqoop that allows you to **retain your job definitions and to easily run them anytime**.

The most important use of the metastore is for running incremental jobs and not having to keep track of the last-value imported.

```
Create job: % sqoop job --create visits -- import --connect ..
```

```
Execute job: % sqoop job --exec visits
```

```
List jobs defined in metastore: % sqoop job --list
```

```
Show saved job configuration: % sqoop job --show visits
```

```
Delete job in metastore: % sqoop job --delete visits
```



FLUME

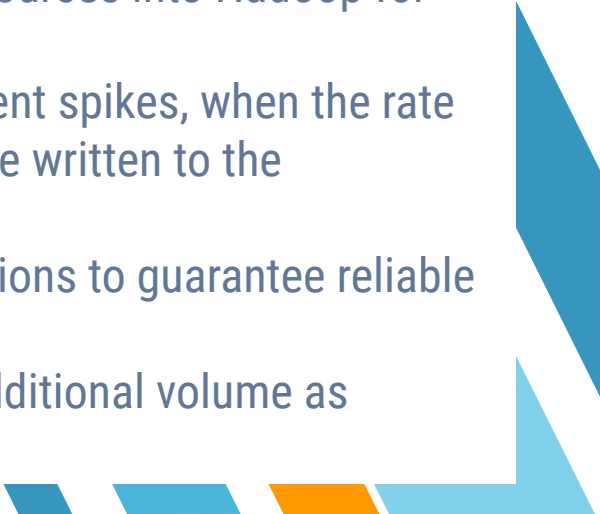
Stream data to HDFS



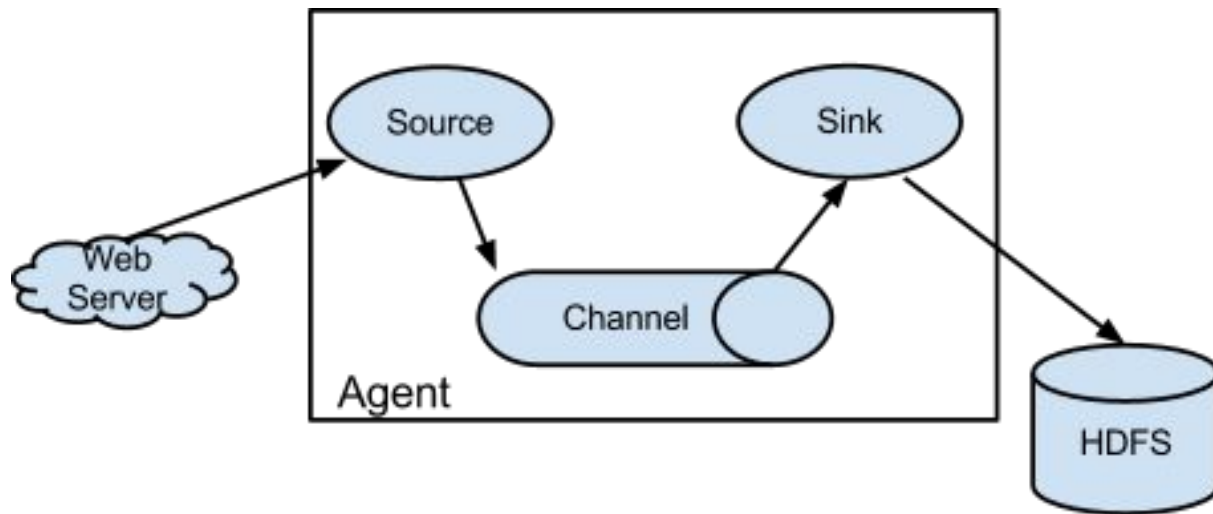
Apache Flume

Apache Flume is a **distributed, reliable, and available system for efficiently collecting, aggregating and moving large amounts of log data** from many different sources to a centralized data store.

Features:


- » **Stream of data:** Ingest streaming data from multiple sources into Hadoop for storage and analysis.
 - » **Insulate systems:** Buffer storage platform from transient spikes, when the rate of incoming data exceeds the rate at which data can be written to the destination.
 - » **Guarantee data delivery:** uses channel-based transactions to guarantee reliable message delivery.
 - » **Scale horizontally:** To ingest new data streams and additional volume as needed.
- 

Apache Flume Agent

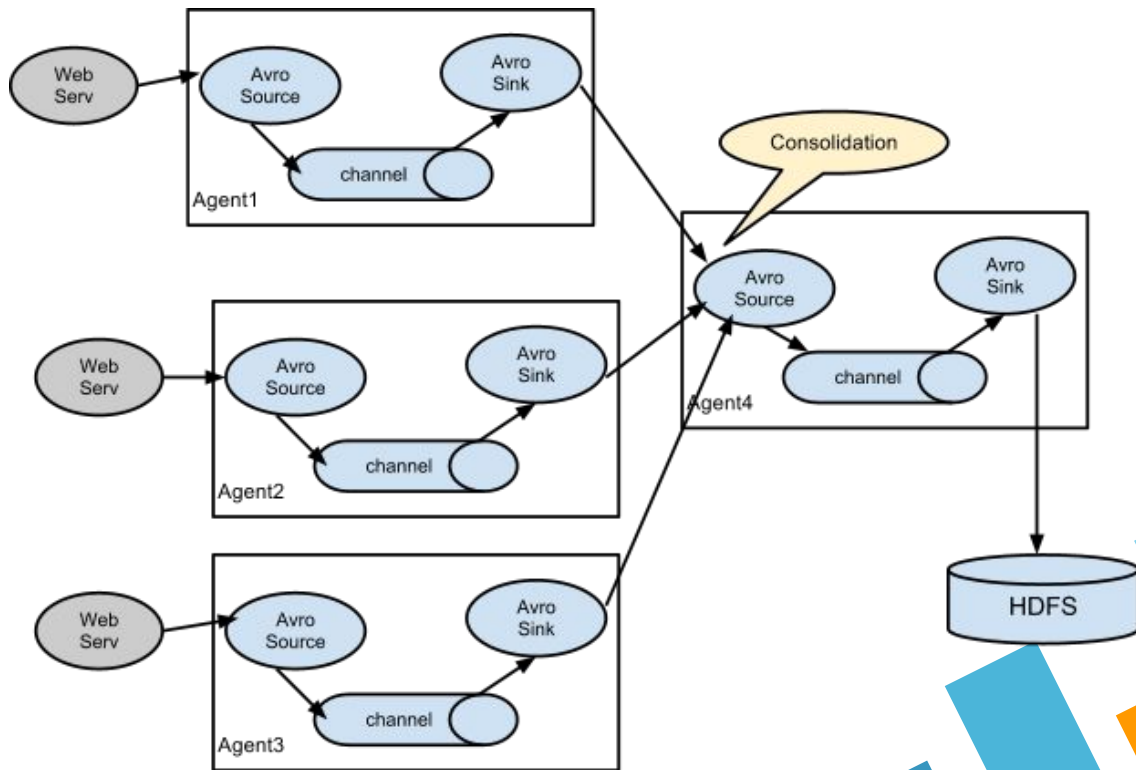




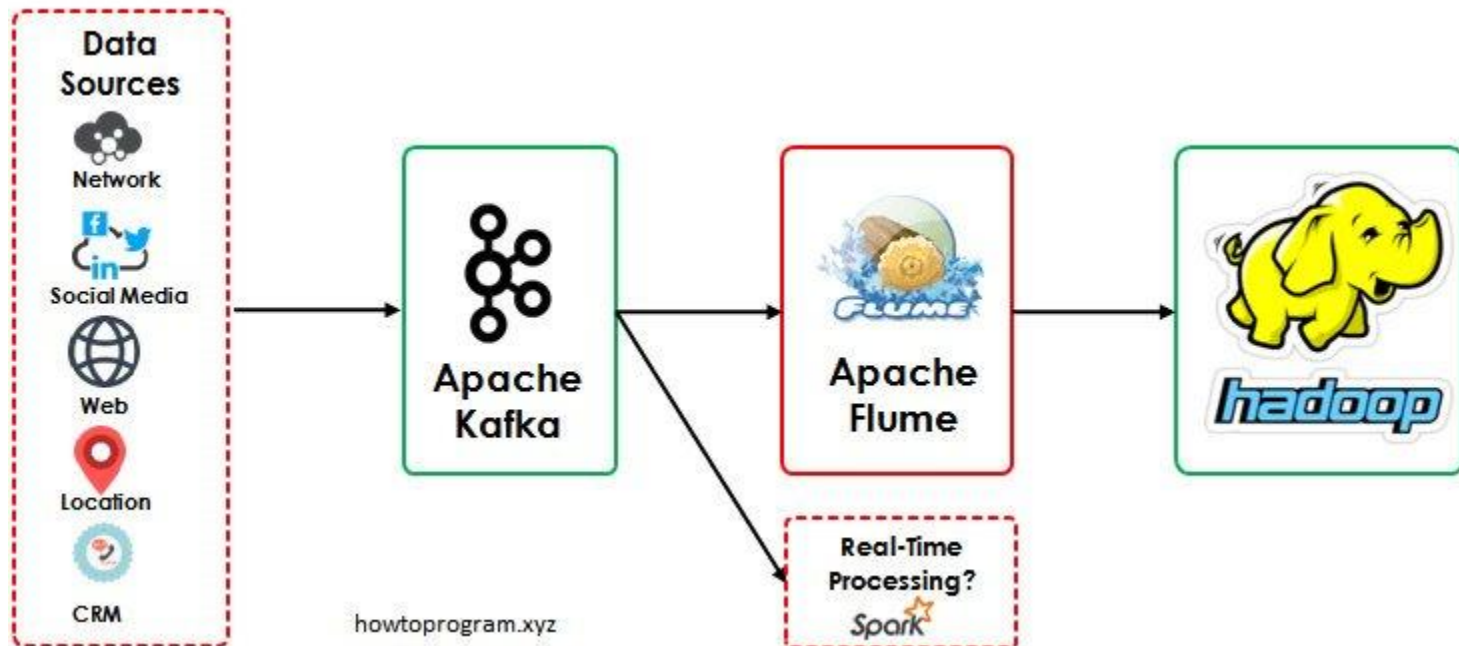
Apache Flume Components

- » **Event:** a singular unit of data that is transported by Flume (typically a single log entry).
 - » **Source:** the entity through which data enters into Flume. Sources either actively poll for data or passively wait for data to be delivered to them. A variety of sources allow data to be collected, such as log4j logs and syslogs.
 - » **Sink:** the entity that delivers the data to the destination. A variety of sinks allow data to be streamed to a range of destinations. One example is the HDFS sink that writes events to HDFS.
 - » **Channel:** the conduit between the Source and the Sink. Sources ingest events into the channel and the sinks drain the channel.
 - » **Agent:** any physical Java virtual machine running Flume. It is a collection of sources, sinks and channels.
 - » **Client:** the entity that produces and transmits the Event to the Source operating within the Agent.
- 

Apache Flume Topology

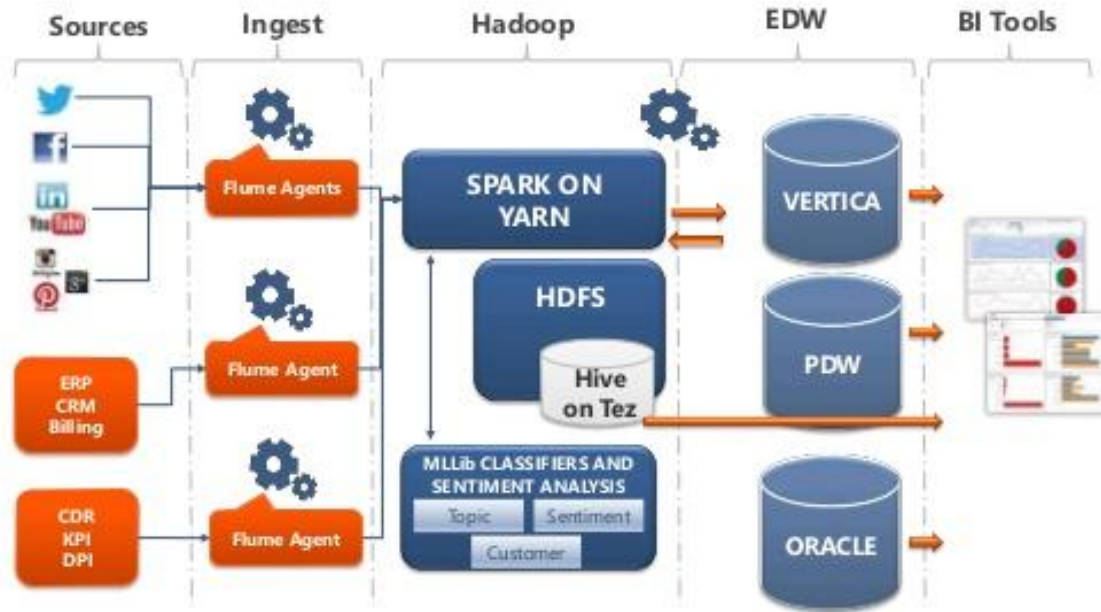


Lambda with Flume



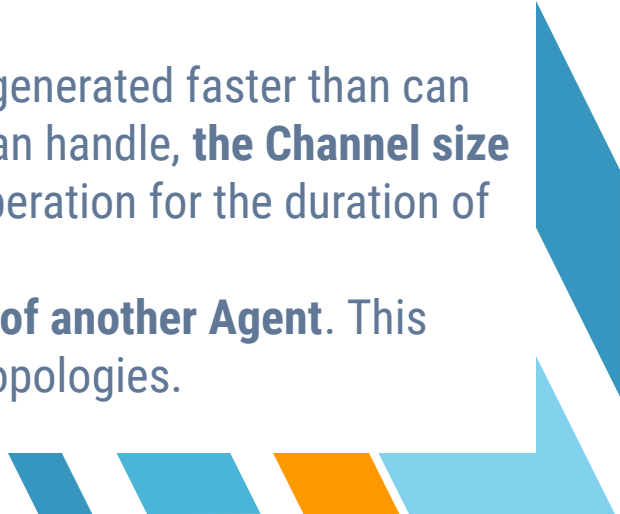
Data Worhouses with Flume

Next Generation Architecture





Apache Flume Components Interactions

1. A **flow** in Flume **starts from the Client**.
 2. The **Client transmits the Event to a Source** operating within the Agent.
 3. The **Source** receiving this Event then **delivers it to one or more Channels**.
 4. One or more **Sinks** operating within the same Agent **drain these Channels**.
 5. **Channels decouple the ingestion rate** from drain rate using the familiar producer-consumer model of data exchange.
 6. **When spikes in client side activity** cause data to be generated faster than can be handled by the provisioned destination capacity can handle, **the Channel size increases**. This allows sources to continue normal operation for the duration of the spike.
 7. **The Sink of one Agent can be chained to the Source of another Agent**. This chaining enables the creation of complex data flow topologies.
- 

Apache Flume Example netcat 2 log

Archivo de configuración (example.conf)

Name the components on this agent

aNet2Log.sources = netcat-source

aNet2Log.channels = memory-channel

aNet2Log.sinks = logger-sink

Describe/configure Source

aNet2Log.sources.netcat-source.type = netcat

aNet2Log.sources.netcat-source.bind = localhost

aNet2Log.sources.netcat-source.port = 44444

Describe the sink

aNet2Log.sinks.logger-sink.type = logger

Use a channel which buffers events in memory

aNet2Log.channels.memory-channel.type = memory

aNet2Log.channels.memory-channel.capacity = 1000

aNet2Log.channels.memory-channel.transactionCapacity =
100

Bind the source and sink to the channel

aNet2Log.sources.netcat-source.channels = memory-channel

aNet2Log.sinks.logger-sink.channel = memory-channel

Para probar (correr en 2 consolas):

```
consola-1> bin/flume-ng agent --name aNet2Log --conf-file  
example.conf -Dflume.root.logger=INFO,console
```

```
consola-2> curl telnet://localhost:44444
```



Apache Flume Example netcat 2 HDFS

Name the components on this agent

aNet2HDFS.sources = netcat-source

aNet2HDFS.channels = memory-channel

aNet2HDFS.sinks = hdfs

Describe/configure Source

aNet2HDFS.sources.netcat-source.type = netcat

aNet2HDFS.sources.netcat-source.bind = localhost

aNet2HDFS.sources.netcat-source.port = 44445

Describe the sink

aNet2HDFS.sinks.hdfs.type = hdfs

aNet2HDFS.sinks.hdfs.hdfs.path = hdfs:/log_data/

aNet2HDFS.sinks.hdfs.hdfs.fileType = DataStream

aNet2HDFS.sinks.hdfs.hdfs.writeFormat = Text

aNet2HDFS.sinks.hdfs.hdfs.batchSize = 1000

aNet2HDFS.sinks.hdfs.hdfs.rollSize = 0

aNet2HDFS.sinks.hdfs.hdfs.rollCount = 10000

Use a channel which buffers events in memory

aNet2HDFS.channels.memory-channel.type = memory

aNet2HDFS.channels.memory-channel.capacity = 1000

aNet2HDFS.channels.memory-channel.transactionCapacity =
100

Bind the source and sink to the channel

aNet2HDFS.sources.netcat-source.channels = memory-channel

aNet2HDFS.sinks.hdfs.channel = memory-channel



Key Take Aways

Tools for ingressing and exporting data to and from HDFS **are fundamental** for any Big Data adoption strategy.

The Hadoop ecosystem offers many different **tools according to the kind of datasource** we are transferring data to and/or from.

It is important to keep in mind that **many datasources are not prepared for the level of parallelization** that Hadoop is able and therefore configurations need to be adjusted accordingly.

The tools described during this presentation are **mostly of batch nature with the exception of Apache Flume** which can handle streaming events of data.





References

- » Mark Grover, Ted Malaska, Jonathan Seidman, Gwen Shapira, **Hadoop Application Architectures**, O'Reilly Media, July 2015
 - » <http://sqoop.apache.org/docs/1.4.5/SqoopUserGuide.html>
 - » Kathleen Ting and Jarek Jarcec Cecho, **Apache Sqoop Cookbook**, O'Reilly Media, July 2013
 - » <http://hortonworks.com/hadoop/flume/>
 - » <https://flume.apache.org/FlumeUserGuide.html>
- 



CREDITS

Content of the slides:

- » Big Data Tools - ITBA

Images:

- » Big Data Tools - ITBA
- » obtained from: commons.wikimedia.org

Special thanks to all the people who made and released these awesome resources for free:

- » Presentation template by [SlidesCarnival](https://slidescarnival.com)
 - » Photographs by [Unsplash](https://unsplash.com)
- 